Nikita Nangia[1]
nikitanangia@nyu.edu

Samuel R. Bowman[1,2,3]
bowman@nyu.edu

[1]Center for Data Science
New York University

[2]Dept. of Linguistics
New York University

[3]Dept. of Computer Science
New York University

Abstract

The GLUE benchmark has seen dramatic progress, with state of the art moving from 69 to over 80 in the past year. Wondering if there is any headroom left, we investigate human performance on GLUE. We attempt to provide a fair (if conservative) estimate of human performance on the benchmark using crowd-worker annotators. Our human performance baseline reaches a score of 86.9, and we see that humans robustly outperform the current state of the art on six of the nine GLUE tasks. Given the fast pace of progress however, the observed headroom is relatively small. We run some experiments in low-resource settings and conclude that working with low-resource datasets could be a valuable next step in natural language understanding research.

1 Introduction

This past year has seen tremendous progress in building general purpose models that can learn good language representations across a range of tasks and domains (Peters et al., 2018; McCann et al., 2017; Devlin et al., 2018). The General Language Understanding Evaluation (GLUE; Wang et al., 2019) benchmark was designed to promote the development of models that can handle many different language understanding tasks and genres without needing to be fully retrained on massive amounts of data. GLUE includes nine different sentence-level natural language understanding (NLU) tasks, including natural language inference, sentiment analysis, acceptability judgment, sentence similarity, and common sense reasoning.

The recent Bidirectional Encoder Representations from Transformers (BERT) model by Devlin et al. (2018) is pre-trained with a language modeling like objective on a large amount of unlabeled data, and then fine-tuned to a specific task.

BERT represents state-of-the-art on GLUE, with a wide margin between it and the next best system (GPT on STILTs; Phang et al., 2019; Radford et al., 2018). BERT's performance on GLUE is impressive enough to now prompt the question: How much better are humans at these NLP tasks? Have modern methods exhausted the headroom in typical sentence-level NLU tasks? In the case of some language understanding tasks, like SQuAD 2.0 (Rajpurkar et al., 2018), the current state-of-the-art model, which is built on top of BERT,[1] is extremely close behind human performance baseline. On the Situations With Adversarial Generations (SWAG; Zellers et al., 2018) dataset, BERT *outperforms* expert human annotators. In this work, we estimate human performance on the GLUE test set to see which tasks have substantial remaining headroom between human and machine performance. We also present some analysis and discussion on what direction we think NLU tasks could go in next.

While human performance or interannotator agreement numbers have been reported on some GLUE tasks, the data collection methods used to establish those baselines vary substantially. To maintain consistency in our reported baseline numbers and to ensure that our results are at least roughly comparable to numbers for submitted machine learning models, we collect annotations using a uniform method for all nine GLUE tasks.

To estimate human performance on GLUE, we conduct a data collection effort with crowdworkers: For each of the nine GLUE tasks, we give the workers a brief training exercise on the task, ask them to annotate a random subset of the evaluation data, and then collect *majority vote* labels from five annotators for each example in the subset. Comparing these labels with the ground-

---

[1]https://rajpurkar.github.io/
SQuAD-explorer/

| | Avg | **Single Sentence** | | **Sentence Similarity** | | | **Natural Language Inference** | | | |
		CoLA	**SST-2**	**MRPC**	**STS-B**	**QQP**	**MNLI**	**QNLI**	**RTE**	**WNLI**
Training Size		*8.5k*	*67k*	*3.7k*	*7k*	*364k*	*393k*	*108k*	*2.5k*	*634*
Human	**86.9**	**66.4**	**97.8**	80.8/86.3	**92.7/92.6**	80.4/59.5	**92.0/92.8**	**91.2**	**93.6**	**95.9**
BERT	80.3	60.5	94.9	**85.4/89.3**	87.6/86.5	**89.3/72.1**	86.7/85.9	**91.1**	70.1	65.1
Δ	6.6	5.9	2.9	-4.6/-3.0	5.1/6.1	-8.9/-12.6	4.1/5.5	0.1	23.5	30.8
BERT-5000	75.8	57.6	92.0	85.4/89.3	87.1/85.8	82.2/61.0	76.4/76.9	89.2	69.2	65.1
BERT-1000	70.7	49.0	90.4	78.5/84.3	83.6/82.3	77.8/55.8	66.5/68.3	86.6	65.6	65.1
BERT-500	68.5	37.2	88.1	74.0/80.7	77.3/75.2	75.4/51.2	61.8/63.0	85.7	61.5	65.1

Table 1: The Human baseline numbers are estimated using no more than 500 test examples. All the BERT scores we report are for BERT-Large. As in the original GLUE paper, we report the Matthews correlation coefficient for CoLA. For MRPC and Quora, we report accuracy and then F1. For STS-B, we report Pearson and then Spearman correlation coefficients. For MNLI, we report accuracy on the matched and then mismatched test sets. For all other tasks we report accuracy. The *Avg* column shows the overall GLUE score: an average across each row, weighting each task equally. The Δ columns shows the difference between the *Human* performance baseline and BERT. The *Training Size* row gives the size of the full training dataset for each task. The BERT-5000/1000/500 rows show test set results for BERT when it is trained on only 5k, 1k, and 500 examples respectively.

truth GLUE test labels yields an overall estimated GLUE score of 86.9—well above BERT's 80.3— and yields single-task scores that are substantially better than BERT on six of nine tasks. However, in light of the progress made on GLUE this past year, the gap in most tasks is relatively small.

The one striking exception is the Winograd Schema NLI Corpus (WNLI; based on Levesque et al., 2012). On this data-poor common-sense reasoning task, humans reach 95.9% accuracy, while no existing machine learning system exceeds the majority-class baseline of 65.1%.

To study BERT's performance on the other GLUE tasks in low-resource settings, we train BERT-Large on just 500, 1000, and 5000 examples. We indeed find that in data-poor versions of the same tasks, BERT suffers considerably.

Ultimately, given the generally impressive performance of BERT on GLUE, we believe that for the future we need tasks that challenge machine learning systems in different ways than our current benchmark tasks. One potential direction is to do more work in data-poor settings to build systems with lower sample complexity.

2 Background and Related Work

GLUE GLUE (Wang et al., 2019) is composed of nine sentence or sentence-pair classification or regression tasks: MultiNLI (Williams et al., 2018), RTE (competition releases 1–3 and 5, merged and treated as a single binary classification task; Dagan et al. 2006, Bar Haim et al. 2006, Giampiccolo et al. 2007, Bentivogli et al. 2009), QNLI (an answer sentence selection task based on SQuAD;

(Rajpurkar et al., 2016)), and WNLI test natural language inference. WNLI is derived from private data created for the Winograd Schema Challenge (Levesque et al., 2012) and it specifically tests for common sense reasoning. The Microsoft Research Paraphrase Corpus (MRPC; Dolan and Brockett, 2005), the Semantic Textual Similarity Benchmark (STS-B; Cer et al., 2017), and Quora Question Pairs (QQP)[2] test paraphrase and sentence similarity evaluation. The Corpus of Linguistic Acceptability (CoLA; Warstadt et al., 2018) tests grammatical acceptability judgment. And lastly, the Stanford Sentiment Treebank (SST; Socher et al., 2013) tests sentiment analysis.

Human Evaluations on GLUE Tasks Warstadt et al. (2018) report human performance numbers on CoLA as well. Using the majority decision from five annotators on 200 examples, they get a Matthews correlation coefficient (MCC) of 71.3. Bender (2015) also estimates human performance on the Winograd Schema Challenge (WSC). They use crowdworkers through Amazon's Mechicanical Turk and report an average accuracy of 92.1%. While they report on the standard WSC, our experiments are on WNLI. Wang et al. (2019) report human performance numbers on GLUE's manually curated diagnostic test set. The examples in this test set are natural language inference sentence pairs that are tagged for a set of linguistic phenomena. They use expert annotators and report an average R_3 coefficient of 0.8.

---

[2] https://data.quora.com/
First-Quora-Dataset-Release-Question-Pairs

To establish human performance on GLUE tasks, we hire annotators through the Hybrid[3] data collection platform, which is similar to Amazon's Mechanical Turk. We conduct the data collection in two phases: Each worker first completes a short training procedure then moves on to the main annotation task. For the main annotation, we tune the pay rate for each task, yielding an average rate of $17/hour. The training phase is short and has a lower, standard pay rate per response, with an average pay of $7.6/hour.

3 Data Collection Method

Training   In the training phase for each GLUE task, each worker answers 20 randomly sampled examples from the task development set. On the training page they are linked to instructions that are tailored to each task. On each page of the worker training, five examples are shown and the answers can be revealed by clicking on a "Show" button at the bottom of the page. The workers are instructed to answer each set of questions and check their work so they can familiarize themselves with the task. Workers who get less than 65% of the examples correct during training are do not qualify for the main task. This is an intentionally low threshold meant only to encourage a reasonable effort. Within our data-collection framework, we cannot fully prevent workers from changing their answers after viewing the correct labels, so we can not use the training phase a as a substantial filter. (See Appendix A.1 for details on the training phase.)

Annotation   Upon finishing training, the workers move onto the annotation phase, which is our source for human performance baseline on GLUE. We randomly sample 500 examples from each task's test set for annotation, with the exception of WNLI where we sample 145 of the 147 available test examples (the two missing examples are the result of a data preparation error). For each of these sampled data points, we collect five annotations from five different workers (see Appendix A.2). We use the test set since the test and development sets are qualitatively different for some tasks, and because we wish compare our results directly with those on the GLUE leaderboard.[4]

4 Results

To calculate the human performance baseline, we take the majority vote across the 5 crowd-sourced annotations. In the case of MultiNLI, since there are three possible labels—*entailment*, *neutral*, and *contradiction*—about 2% of examples see a tie between two labels. For these ties, we take the label that is more frequent in the development set. In the case of STS-B, we take an average of the scalar annotator labels. Since we only collect annotations for a subset of the data, we could not access the test set through the GLUE leaderboard interface. Instead, we worked in cooperation with the GLUE team to measure performance.

Human performance appears in the first line of Table 1. These results show that our annotators outperform BERT overall on GLUE. The human baseline beats GLUE on six of the single-tasks, however the margin is not considerable in five of them. On MRPC and QQP, the BERT machine *outperforms* our annotators by a sizeable margin. The results on QQP are particularly surprising: BERT scores 12.6 F1 points better than our annotators. Our annotators however, are only given 20 examples and a short set of instructions to train them on GLUE tasks. By comparison, BERT is fine-tuned on the full 364k-example QQP training set. The discrepancy in the amount of training data may be particularly pertinent for paraphrase tasks because the task is a little subjective. For example, the following pairs from QQP's development set are labeled as a duplicates[5],

> Question-1: *"What is actual meaning of life? Indeen, it depend on perception of people or other thing?"*; Question-2: *"What is the meaning of my life?"*

> Question-1: *"How do you know if you're in love?"*; Question-2: *"How can you know if you're in love or just attracted to someone?"*

In both pairs, one of the questions asks a more detailed, specific question. A reasonable reader could interpret these questions as asking different things. If given more training data, it is possible that our annotators could better learn the peculiar labeling definitions fitting the QQP corpus.

BERT's reliance on large training data may be further evidenced by its performance discrepancy between MultiNLI and RTE: human performance is quite similar for the two, but BERT does over 15 points better on MultiNLI. Both MultiNLI and RTE are textual entailment datasets, but MultiNLI's training set is quite large at 393k examples, while the GLUE version of RTE has only 2.5k examples.

To better investigate BERT's sample complexity, we train it on 5k, 1k, and 500 examples for each GLUE task (or fewer for tasks with fewer training examples). We use the publicly available implementation of BERT-Large released by Devlin et al. (2018). We use the publicly distributed pretrained weights as the initialization for fine-tuning on the GLUE tasks. We also use the hyperparameters reported by Devlin et al. (2018). The results for these experiments are shown in the last three lines of Table 1. We see a precipitous drop in performance on most tasks with large datasets. One exception here is QNLI. One possible explanation is that both the QNLI source texts and the BERT trainining data come from English Wikipedia (Rajpurkar et al., 2016; Wang et al., 2019). On MRPC and QQP however, BERT's performance drops below human performance in the 1k and 500-example settings.

We would like to note that our human performance number on CoLA is 4.9 points below what was reported in Warstadt et al. (2018). We believe this discrepancy is because they use Linguistics PhD students as expert annotators while we use crowdworkers. This further supports our belief that our human performance baseline is a conservative estimate, and that higher performance is possible, particularly with more training.

5   Discussion

Our estimate of human performance shows that human annotators can beat the state-of-the-art BERT system on GLUE by at least 6.6 points. While the human baseline is better than BERT on six of nine individual tasks, the rapid rate of progress on GLUE indicates that the margins on these tasks is relatively small. The one exception here is WNLI.

No system on the GLUE leaderboard has managed to exceed the performance of the most-frequent-class baseline on WNLI, and several papers that propose methods for GLUE justify their poor performance by asserting that the task must be somewhat broken.[6] While WNLI was constructed so as not to include any statistical cues that a simple machine learning system can exploit, which can make it quite difficult, the WNLI test set nonetheless shows one of the *highest* human performance scores of the nine GLUE tasks, reflecting its status as a corpus constructed and vetted by artificial intelligence experts. This makes it clear that tasks like WNLI with small training sets (634 sentence pairs) and no simple cues remain a serious (and sometimes unacknowledged) blind spot for modern neural network sentence understanding methods.

We do find that BERT outperforms the human baseline on MRPC and QQP. A qualitative analysis of the examples that our annotators get wrong on MRPC and QQP shows that the labels in these instances do not not often match the colloquial meaning of *parapharse* and that the labels are debatable. Given the massive amount of training data, machine learning systems are able to tune themselves to the label definitions and pick up on aggregate statistics that our annotators do not have access to. It is possible that with more training human annotators will be able to match machine performance on these two tasks.

In our data-constrained fine-tuning experiments with BERT, we see that BERT suffers in low-resource settings. This result gives us more reason to believe that low-resource settings continue to be challenging for machine learning systems. If we want more robust, flexible, and easily adaptable machine systems, designing them to have low sample complexity will be a step in the right direction.

6   Conclusion

This paper presents a conservative estimate of human performance to serve as a performance target for the GLUE sentence understanding benchmark. We obtain this baseline with the help of crowdworker annotators. We see that state-of-the-art models like BERT are not far behind human performance on most GLUE tasks, but we also note that, when trained in low-resource settings, BERT's performance falls considerably. Given these results, and the continued difficulty neural methods have with the Winograd Schema Chal-

---

[6]Devlin et al. (2018), for example, mention that they avoid "the problematic WNLI set".

lenge, we propose that future NLU benchmark datasets could provide valuable challenges by having smaller training set.

Acknowledgments

References

Roy Bar Haim, Ido Dagan, Bill Dolan, Ferro Lisa, Danilo Giampiccolo, Bernardo Magnini, and Idan Szpektor. 2006. The second PASCAL recognising textual entailment challenge. In *Proceedings of the Second PASCAL Challenges Workshop on Recognising Textual Entailment*.

David Bender. 2015. Establishing a human baseline for the winograd schema challenge. In *MAICS*.

Luisa Bentivogli, Ido Dagan, Hao Trang Dang, Danilo Giampiccolo, and Bernardo Magnini. 2009. The fifth PASCAL recognizing textual entailment challenge. In *Proceedings of Text Analysis Conference*.

Daniel Cer, Mona Diab, Eneko Agirre, Inigo Lopez-Gazpio, and Lucia Specia. 2017. Semeval-2017 task 1: Semantic textual similarity multilingual and crosslingual focused evaluation. In *Proceedings of the 11th International Workshop on Semantic Evaluation (SemEval-2017)*, pages 1–14. Association for Computational Linguistics.

Ido Dagan, Oren Glickmen, and Magnini Bernardo. 2006. The PASCAL recognising textual entailment challenge. In *Machine learning challenges*, pages 177–190. Springer.

Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. BERT: pre-training of deep bidirectional transformers for language understanding. *arXiv preprint 1810.04805*.

William B Dolan and Chris Brockett. 2005. Automatically constructing a corpus of sentential paraphrases. In *In Proceedings of IWP*.

Danilo Giampiccolo, Bernardo Magnini, Ido Dagan, and Bill Dolan. 2007. The third PASCAL recognizing textual entailment challenge. In *ACL-PASCAL workshop on textual entailmentand paraphrasing*, pages 1–9. Association for Computational Linguistics.

Hector J. Levesque, Ernest Davis, and Leora Morgenstern. 2012. The winograd schema challenge. In *Proceedings of the Thirteenth International Conference on Principles of Knowledge Representation and Reasoning*, KR'12, pages 552–561. AAAI Press.

Bryan McCann, James Bradbury, Caiming Xiong, and Richard Socher. 2017. Learned in translation: Contextualized word vectors. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems 30*, pages 6294–6305. Curran Associates, Inc.

Matthew Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. 2018. Deep contextualized word representations. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Association for Computational Linguistics.

Jason Phang, Thibault Févry, and Samuel R. Bowman. 2019. Sentence encoders on stilts: Supplementary training on intermediate labeled-data tasks. *arXiv preprint 1811.01088*.

Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. 2018. Improving language understanding by generative pre-training. Unpublished ms. available through a link at https://blog.openai.com/language-unsupervised/.

Pranav Rajpurkar, Robin Jia, and Percy Liang. 2018. Know what you don't know: Unanswerable questions for SQuAD. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 784–789. Association for Computational Linguistics.

Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392. Association for Computational Linguistics.

Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng, and Christopher Potts. 2013. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics.

Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R. Bowman. 2019. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *International Conference on Learning Representations*.

Alex Warstadt, Amanpreet Singh, and Samuel R. Bowman. 2018. Neural network acceptability judgments. *arXiv preprint 1805.12471*.

Adina Williams, Nikita Nangia, and Samuel Bowman. 2018. A broad-coverage challenge corpus for sentence understanding through inference. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1112–1122. Association for Computational Linguistics.

Rowan Zellers, Yonatan Bisk, Roy Schwartz, and Yejin Choi. 2018. Swag: A large-scale adversarial dataset for grounded commonsense inference. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics.

A Crowd-Sourced Data Collection

A.1 Training Phase

During training, we provide a link to task-specific instructions. As an example, the instructions for CoLA are shown in Table 2. The instructions for all tasks follow the same format: briefly describing the annotator's job, explaining the labels, and providing at least one example. All of the instruction and FAQ pages are publicly available[7].

In addition to the task-specific instruction, we also provide general instructions on the training phase. An example of these instructions is shown in Table 3. The only variation from task to task, is the name of task in the instructions. Lastly, we provide a link to an FAQ page. The FAQ page addresses the balance of the data. If the labels are balanced, we tell the annotators so. If the labels are not balanced, we assure the annotators that they need not worry about assigning one label more frequently. For most tasks we also inform the annotators where the data comes from, for example from another crowdsourcing effort or from news articles.

During training, each page is headed by the instructions and then the annotator is given five examples to label. At the bottom of the page, there is a "Show" button which reveals the answers. If their submitted answer was incorrect, the correct label is shown in red, otherwise it's in black. In the instructions, the worker is asked to check their work with this button.

A.2 Annotation Phase

In the main data collection phase we simply provide the annotator a link to the same task-specific

---

[7] https://nyu-mll.github.io/GLUE-human-performance/

instructions used in the training phase (Figure 2). We also provide a link to the same FAQ page as in the training phase. We enforce the training phase as a qualification for annotation, so crowdworkers can not participate in annotation without first completing the associated training.

B QQP Example

The 25 randomly sampled examples from the QQP development set are as follows,

1. Question 1: *"What is actual meaning of life? Indeen, it depend on perception of people or other thing?"*
   Questions 2: *"What is the meaning of my life?"*
   Label: 1

2. Question 1: *"What is the difference between CC and 2S classes of travel in Jan Shatabdi express?"*
   Questions 2: *"What is TQWL in IRCTC wait list? "*
   Label: 0

3. Question 1: *"What would have happened if Hitler hadn't declared war on the United States after Pearl Harbor?"*
   Questions 2: *"What would have happened if the United States split in two after the revolutionary war?"*
   Label: 0

4. Question 1: *"What should you do if you meet an alien?"*
   Questions 2: *"What could be the possible conversation between humans and aliens on their first meeting?"*
   Label: 0

5. Question 1: *"Which European countries provide mostly free university education to Indian citizen? "*
   Questions 2: *"What countries provide free education to Indian students?"*
   Label: 0

6. Question 1: *"What is the meaning of life? Whats our purpose on Earth?"*
   Questions 2: *"What actually is the purpose of life?"*
   Label: 1

Table 2: The instructions given to crowd-sourced worker for the CoLA task. While the instructions were tailored for each task in GLUE, they all followed a similar format.

Table 3: Instructions about the training phase provided to workers. This example is for CoLA training. The only change in instructions for other tasks is the name of the task.

7. Question 1: *"I got 112 rank in CDAC (A+B+C). My subject of interest is VLSI. Is there any chance that I would get CDAC Pune, Noida for VLSI?"*
Questions 2: *"Suggest some good indian youtube channels for studying Aptitude?"*
Label: 0

8. Question 1: *"What are the positives and negatives of restorative justice?"*
Questions 2: *"Is Vengence and Justice opposite?"*
Label: 0

9. Question 1: *"What's a good way to make money through effort? "*
Questions 2: *"How do I make money without much effort?"*
Label: 0

10. Question 1: *"Why can't I ask any questions on Quora?"*
Questions 2: *"Can you ask any question on Quora?"*
Label: 0

11. Question 1: *"Will it be a problem if a friend deposits 4 lakhs in my savings bank account and I don't have a source of income to show?"*
Questions 2: *"I am 25.5 year old boy with a B.Com in a sales job having a package of 4 LPA. I will be married in less than a year. I want to quit my job and start my own business with the savings I have of 2 Lakh. Is this an ideal situation to take a risk?"*
Label: 0

12. Question 1: *"Which among five seasons (summer, winter, autumn, spring, rainy) is*

most favourable for farming and cultivating of crops?”
Questions 2: *“Which among the five seasons (summer, winter, rainy, spring, autumn) is better for farming and cultivating of crops?”*
Label: 1

13. Question 1: *“Should I move from the USA to India?”*
Questions 2: *“Moving from usA to India?”*
Label: 1

14. Question 1: *“How can I find the real true purpose of my life?”*
Questions 2: *“What should one do to find purpose of one's life?”*
Label: 1

15. Question 1: *“Is Donald Trump likely to win the 2016 election (late 2015 / early 2016)?”*
Questions 2: *“What will Donald Trump's response be if he doesn't win the 2016 presidential election?”*
Label: 0

16. Question 1: *“What is the easiest and cheapest way to lose weight fast?”*
Questions 2: *“What are the easiest and the fastest ways to lose weight?”*
Label: 1

17. Question 1: *“Why are basically all of my questions on Quora marked as 'needing improvement'? Am I that bad?”*
Questions 2: *“Why do questions get marked for 'needing improvment' when they clearly don't? ”*
Label: 1

18. Question 1: *“What are some of the most visually stunning apps?”*
Questions 2: *“What are the most visually stunning foods?”*
Label: 0

19. Question 1: *“What are some of the good hotels near chennai central railway station?”*
Questions 2: *“Best places to eat in Chennai?”*
Label: 0

20. Question 1: *“How do you prepare for a job interview?”*
Questions 2: *“How do I prepare for my first job interview?”*
Label: 1

21. Question 1: *“What are the best resources for learning Ukrainian?”*
Questions 2: *“What are the best resources for learning Turkish?”*
Label: 0

22. Question 1: *“How much time will it take to charge a 10,000 mAh power bank?”*
Questions 2: *“How much time does it takes to charge the power bank 13000mAh for full charge?”*
Label: 0

23. Question 1: *“How do you know if you're in love?”*
Questions 2: *“How can you know if you're in love or just attracted to someone?”*
Label: 1

24. Question 1: *“Which are the best and affordable resorts in Goa?”*
Questions 2: *“What are some affordable and safe beach resorts in Goa?”*
Label: 1

25. Question 1: *“How winning money from YouTube?”*
Questions 2: *“How do I make money from a YouTube channel?”*
Label: 1